

AI/ML Engineer - Senior Level

About the Company:

TechCorp is a leading technology company focused on AI-driven solutions for enterprise customers. We're looking for experienced AI/ML engineers to join our growing team.

Job Description:

We are seeking a Senior AI/ML Engineer to design, develop, and deploy machine learning models and systems. The ideal candidate will have strong experience in building production-ready ML pipelines and working with large-scale data.

Key Responsibilities:

- Design and implement end-to-end machine learning pipelines for various business applications
- Develop and deploy NLP models using transformer architectures and BERT-based models
- Build RAG (Retrieval-Augmented Generation) systems for intelligent document processing
- Architect scalable ML infrastructure on cloud platforms (AWS, GCP, Azure)
- Collaborate with cross-functional teams to integrate ML solutions into production systems
- Optimize model performance and establish ML best practices

Technical Requirements:

- 4+ years of experience in machine learning and AI development
- Strong proficiency in Python and ML frameworks (TensorFlow, PyTorch, Scikit-learn)
- Experience with NLP and transformer models (BERT, GPT, etc.)
- Hands-on experience with RAG systems and vector databases
- Expertise in cloud platforms (AWS, Google Cloud, or Azure)
- Experience with MLOps tools and model deployment
- Knowledge of distributed computing and big data technologies
- Familiarity with containerization (Docker, Kubernetes)

Nice to Have:

- Experience with LangChain or similar ML orchestration frameworks
- Background in full-stack development
- Research experience in ML/AI
- Contributions to open-source ML projects

Benefits:

- Competitive salary and equity package
- Health, dental, and vision insurance
- Flexible work arrangements
- Professional development budget
- Collaborative and innovative work environment

Location: San Francisco, CA (Hybrid)